

14 September 2020

ML Benchmarks on Discovery Cluster

Research Computing

Strictly private and confidential



Table of contents

1	Configuration	3
---	---------------	---

01

Configuration

Ensuring **exclusive** GPU usage

Non Exclusive Request for NVIDIA Tesla P100 GPU

```
$ srun --partition gpu --pty --gres=gpu:p100 /bin/bash
```

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
9116	e.dorari	20	0	15.2g	845716	114408	R	99.7	0.2	102:16.46	G82D-alpha+
17987	s.chakr+	20	0	162368	2524	1584	R	0.7	0.0	0:00.09	top
11059	root	0	-20	15.4g	1.3g	112952	S	0.3	0.3	39:19.04	nmfsd
17939	root	20	0	308372	4620	3304	S	0.3	0.0	0:00.07	slurmstepd
1	root	20	0	51976	4148	2620	S	0.0	0.0	2:44.72	systemd

Discovery user **e.dorari** can be seen hogging **99.7%** of the available E5-2680v4@2.40GHz CPU which would heavily skew results

Non Exclusive Request for NVIDIA V100 SXM2 GPU

```
$ srun --partition gpu --pty --gres=gpu:v100-sxm2 /bin/bash
```

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
243212	jefftian	20	0	7460388	201700	114600	R	99.7	0.1	97:01.82	continuu+
243374	jefftian	20	0	7460388	203472	114600	R	99.7	0.1	96:04.34	continuu+
243199	jefftian	20	0	7472108	209904	114600	R	99.0	0.1	96:56.08	continuu+
9	root	20	0	0	0	0	S	0.3	0.0	32:38.49	rcu_sched
243355	root	20	0	305284	4512	3216	S	0.3	0.0	0:00.28	slurmste+
243377	root	20	0	0	0	0	S	0.3	0.0	0:01.02	nv_queue
251939	s.chakr+	20	0	162412	2584	1584	R	0.3	0.0	0:00.36	top

Discovery user **jeff.tian** can be seen hogging **99.7%** of the available Intel Gold 6132@2.60Ghz CPU which would heavily skew results

Exclusive Request for NVIDIA V100 SXM2 GPU

```
$ srun --partition gpu --pty --gres=gpu:p100 --exclusive /bin/bash
```

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
766	s.chakr+	20	0	162368	2524	1584	R	0.3	0.0	0:00.20	top
1	root	20	0	51976	4152	2620	S	0.0	0.0	2:47.79	systemd
2	root	20	0	0	0	0	S	0.0	0.0	0:00.23	kthreadd
4	root	0	-20	0	0	0	S	0.0	0.0	0:00.00	kworker/0+

My user **s.chakravarty** is now free to utilize all the compute

Exclusive Request for NVIDIA V100 SXM2 GPU

```
$ srun --partition gpu --pty --gres=gpu:v100-sxm2 --exclusive /bin/bash
```

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
246024	s.chakr+	20	0	162376	2532	1584	R	0.7	0.0	0:00.12	top
1	root	20	0	52808	4992	2612	S	0.3	0.0	3:28.23	systemd
29	root	rt	0	0	0	0	S	0.3	0.0	0:03.28	watchdog+
12373	root	0	-20	15.5g	1.4g	219780	S	0.3	0.7	54:34.75	nmfsd

My user **s.chakravarty** is now free to utilize all the compute

Creating a **miniconda** environment

Clean Conda Environment Steps

Download Miniconda 2 from the internet

```
$ wget https://repo.anaconda.com/miniconda/Miniconda2-latest-Linux-x86_64.sh
```

Change the permissions of the installation script

```
$ chmod +x Miniconda2-latest-Linux-x86_64.sh
```

Run the installation script to install Miniconda 2

```
$ ./Miniconda2-latest-Linux-x86_64.sh
```

```
Agree to license agreement >> yes
```

```
Directory to install >> /work/rc/s.chakravarty
```

```
cd /work/rc/s.chakravarty/bin
```

Activate your base miniconda environment

```
$ source activate
```

Update all your conda packages

```
conda update conda
```

```
Proceed? >> yes
```

Configuring Conda Environments



RAPIDS

rapids.ai **cuML**

Steps

```
$ conda create --name cuml_env --no-default-packages
```

```
$ conda activate cuml_env
```

```
(cuml_env) $ conda install -c rapidsai-nightly -c nvidia -c conda-forge -c defaults  
rapids=0.16 python=3.8 cudatoolkit=11.0
```



Intel® **DAAL4py**

Steps

```
$ conda create --name daal4py_env --no-default-packages
```

```
$ conda activate daal4py_env
```

```
(daal4py_env) $ conda install -c intel daal4py
```



H₂O.ai

H₂O.ai **H2O4GPU**

Steps

```
$ conda create --name h2o4gpu_env --no-default-packages
```

```
$ conda activate pydaal_env
```

```
(h2o4gpu_env) $ conda create -n h2o4gpuenv -c h2oai -c conda-forge -c rapidsai  
h2o4gpu-cuda10
```

Configuring Conda Environments



RAPIDS

rapids.ai **cuML**

Sample Code

```
Python 3.8.5 | packaged by conda-forge |
>>> from cuml.cluster import KMeans
>>> import cudf, numpy as np, pandas as pd
>>> def np2cudf(df):
...     df = pd.DataFrame({'fea%d'%i:df[:,i] for i in range(df.shape[1])})
...     pdf = cudf.DataFrame()
...     for c,column in enumerate(df):
...         pdf[str(c)] = df[column]
...     return pdf
>>> kmeans_float = KMeans(n_clusters=2).fit(np2cudf(np.array([[1., 1.], [1., 4.], [1., 0.]])
>>> print(kmeans_float.cluster_centers_)
array([[1. , 0.5], [1. , 4. ]])
```



Intel® **DAAL4py**

Sample Code

```
$ python3
Python 3.7.7 :: Intel(R) Corporation
>>> from daal4py import kmeans_init
>>> import numpy as np
>>> X = np.array([[1., 1.], [1., 4.], [1., 0.]])
>>> kmi = kmeans_init(10, method="plusPlusDense")
>>> result = kmi.compute(X)
>>> print(result.centroids)
array([[1. , 0.5], [1. , 4. ]])
```



H₂O.ai

H₂O.ai **H2O4GPU**

Sample Code

```
$ python3
Python 3.6.11 | packaged by conda-forge |
Type "help", "copyright", "credits" or "license" for more information
>>> import h2o4gpu
>>> import numpy as np
>>> X = np.array([[1., 1.], [1., 4.], [1., 0.]])
>>> model = h2o4gpu.KMeans(n_clusters=2, random_state=1234).fit(X)
>>> model.cluster_centers_
array([[1. , 0.5], [1. , 4. ]])
```